Agenda.

→ Design S3 → AWS.

→ File Storage (for storing big files)

- Blob Storage
- S3
- HDFS.

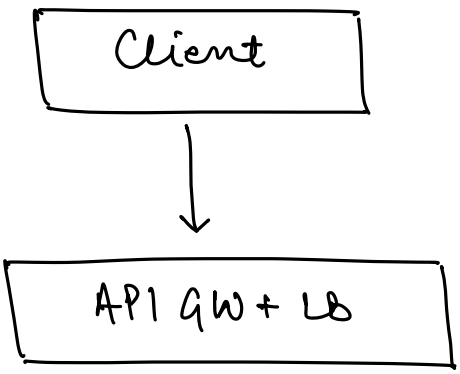

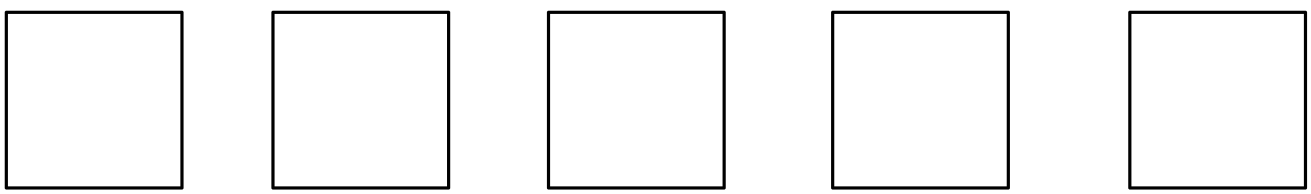

Posts.

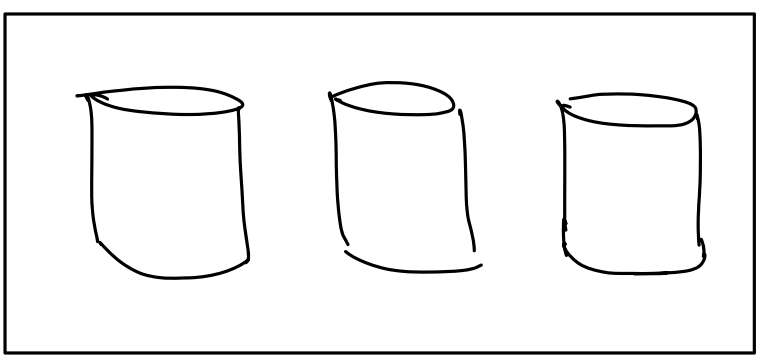

S3 file Storage

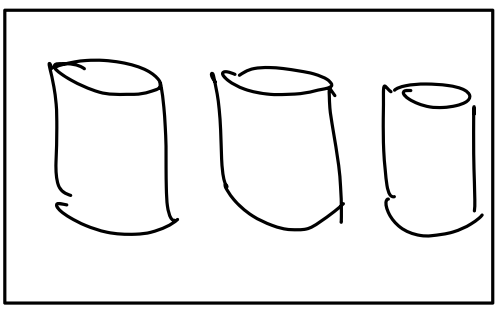

⇒ File Storage.

→ Media Content
(Image | Video | file etc)

→ In main DB (SQL or NOSQL), we just store the S3 URL to refer the content from S3.

→ File Storage System should be able to store big files.

↳ 1GB | 10GB | 1TB | 10TB ...

→ Durable

→ Performance of uploading or downloading.

↳ What happens if our connection drops in b/w ?

Option 1: File stored as 1 unit in 1 machines.

Pros: 1) No need to maintain the chunk info. (Cost of entries)

2) We need not to collate the chunks.

Cons: 1) File size is limited by machine size.

2) Parallelism NOT possible.

Option 2: Divide file into smaller chunks and store each chunk into a different m/c.

Pros:

Cons:

Note: We are going to divide a file into multiple chunks so that the file size isn't limited by the m/c size & parallelism would be possible.

But we shouldn't be dividing file into very small chunks otherwise managing these many chunks would become an overhead.

⇒ HDFS.

↳ Hadoop Distributed file System.

1) Data Nodes : Machines where we store file Chunks. → Repliation

2) Name Nodes : Store metadata about the files.

→ Maintains the mapping of file chunks and in which m/c these chunks are present.

→ Repliation

F1 → 200 Chunks

F1C1 → D7

F1C2 → D10

F1C3 → D4

F1C4 → D8

⋮ ⋮

⇒ In one HDFS Cluster, we don't shard a Name Node server but replication would be there.

⇒ HDFS Cluster.

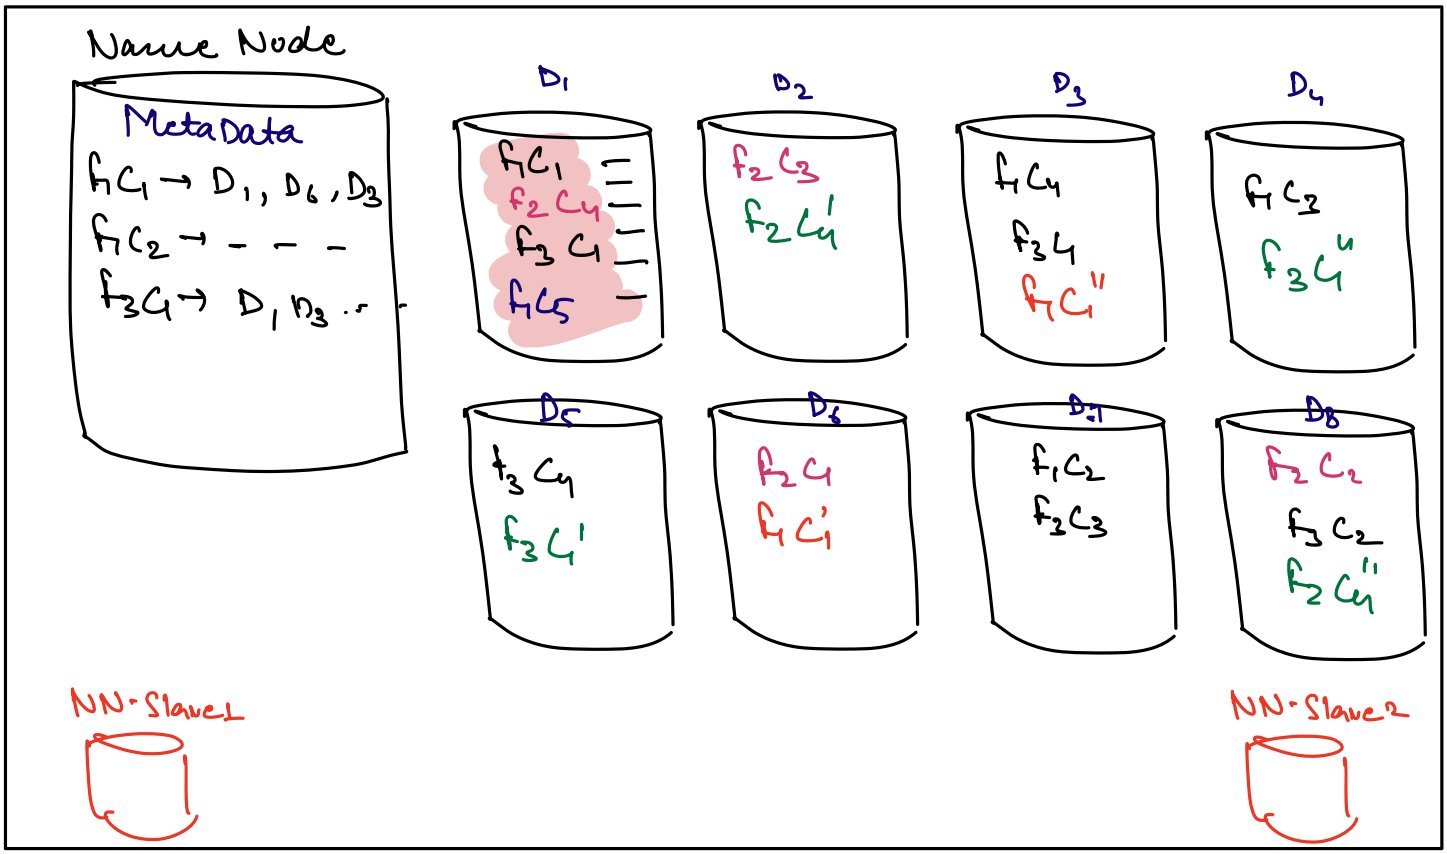

Rack Aware Algo

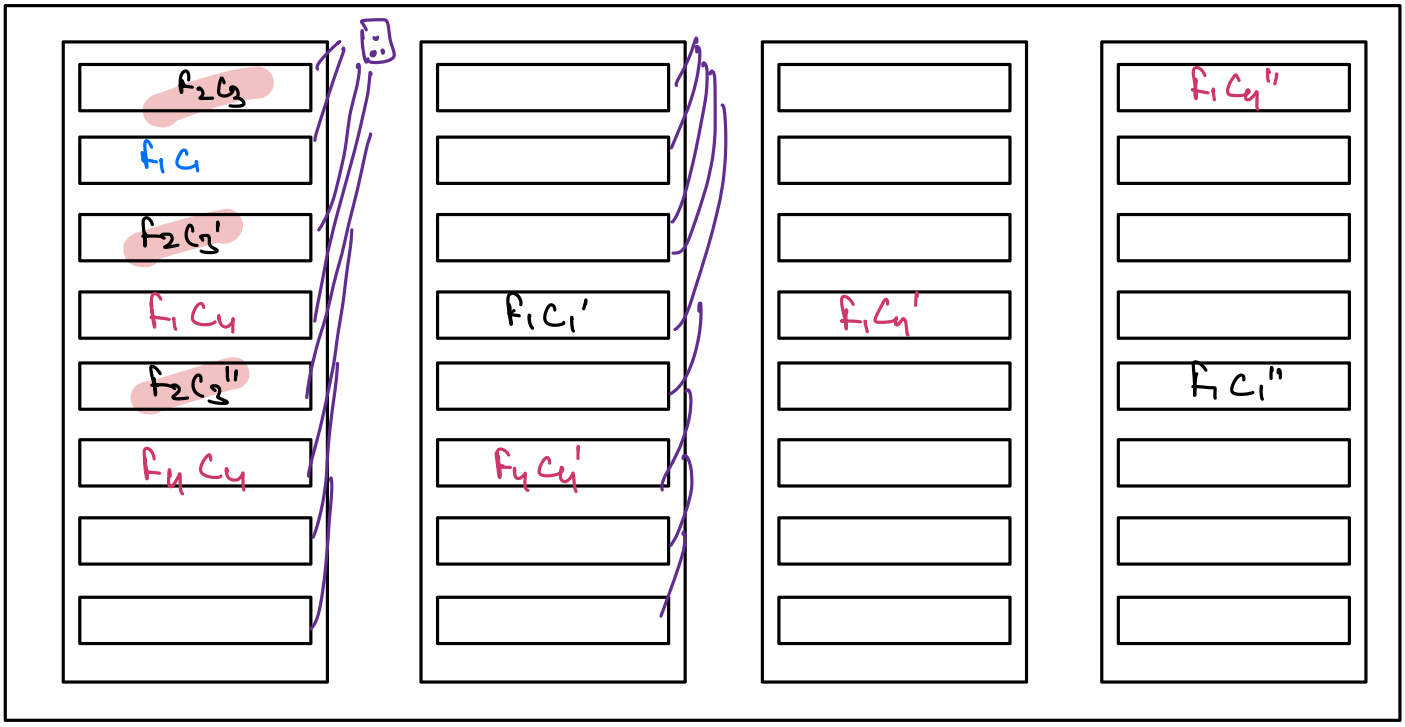

for more reliability, HDFS keeps data on different racks so that we do not loose our data even if a rack goes down.

→ Avoid replicating same chunks within the same rack.

Upload flow.

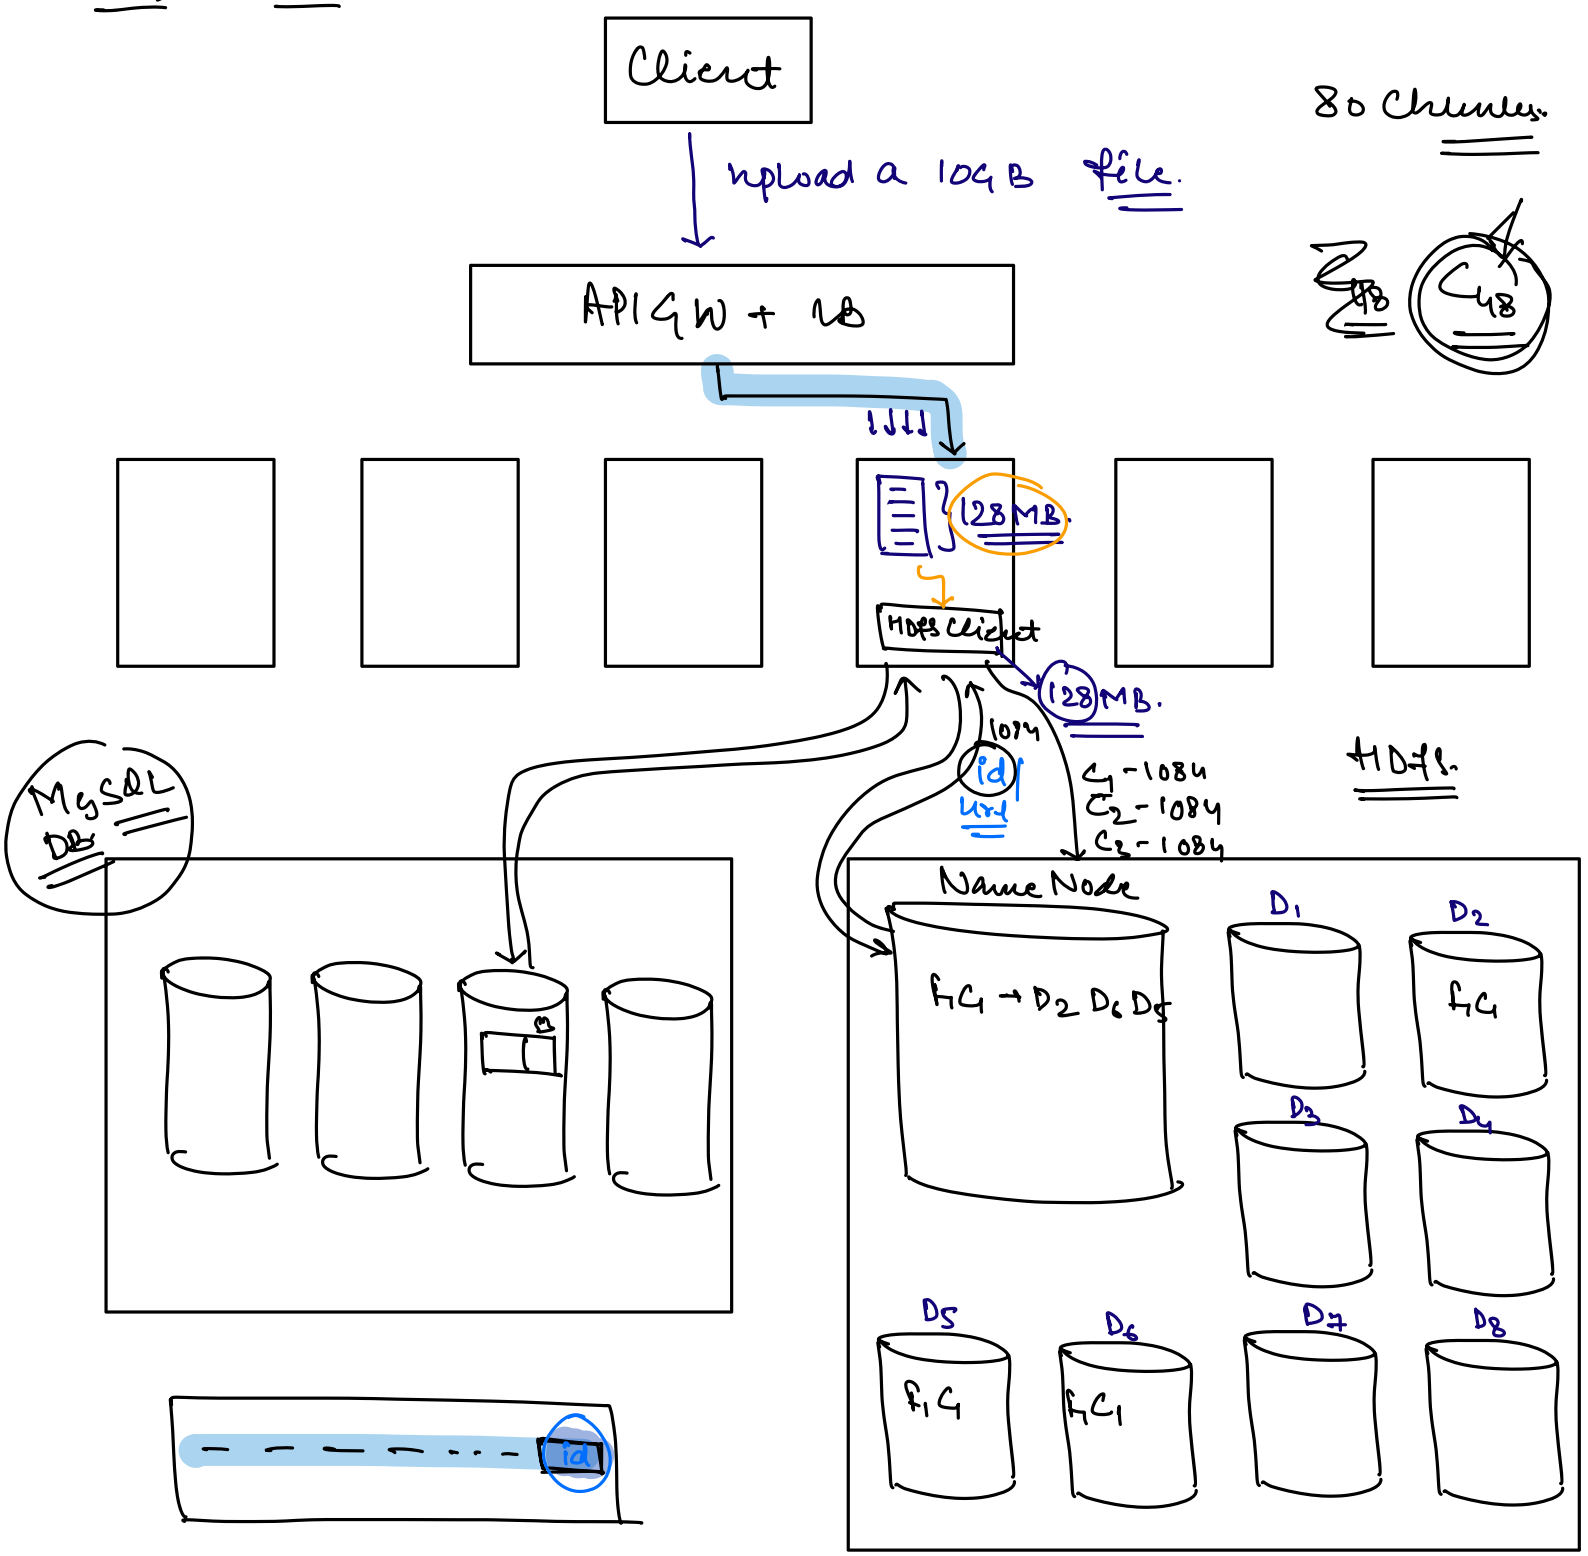

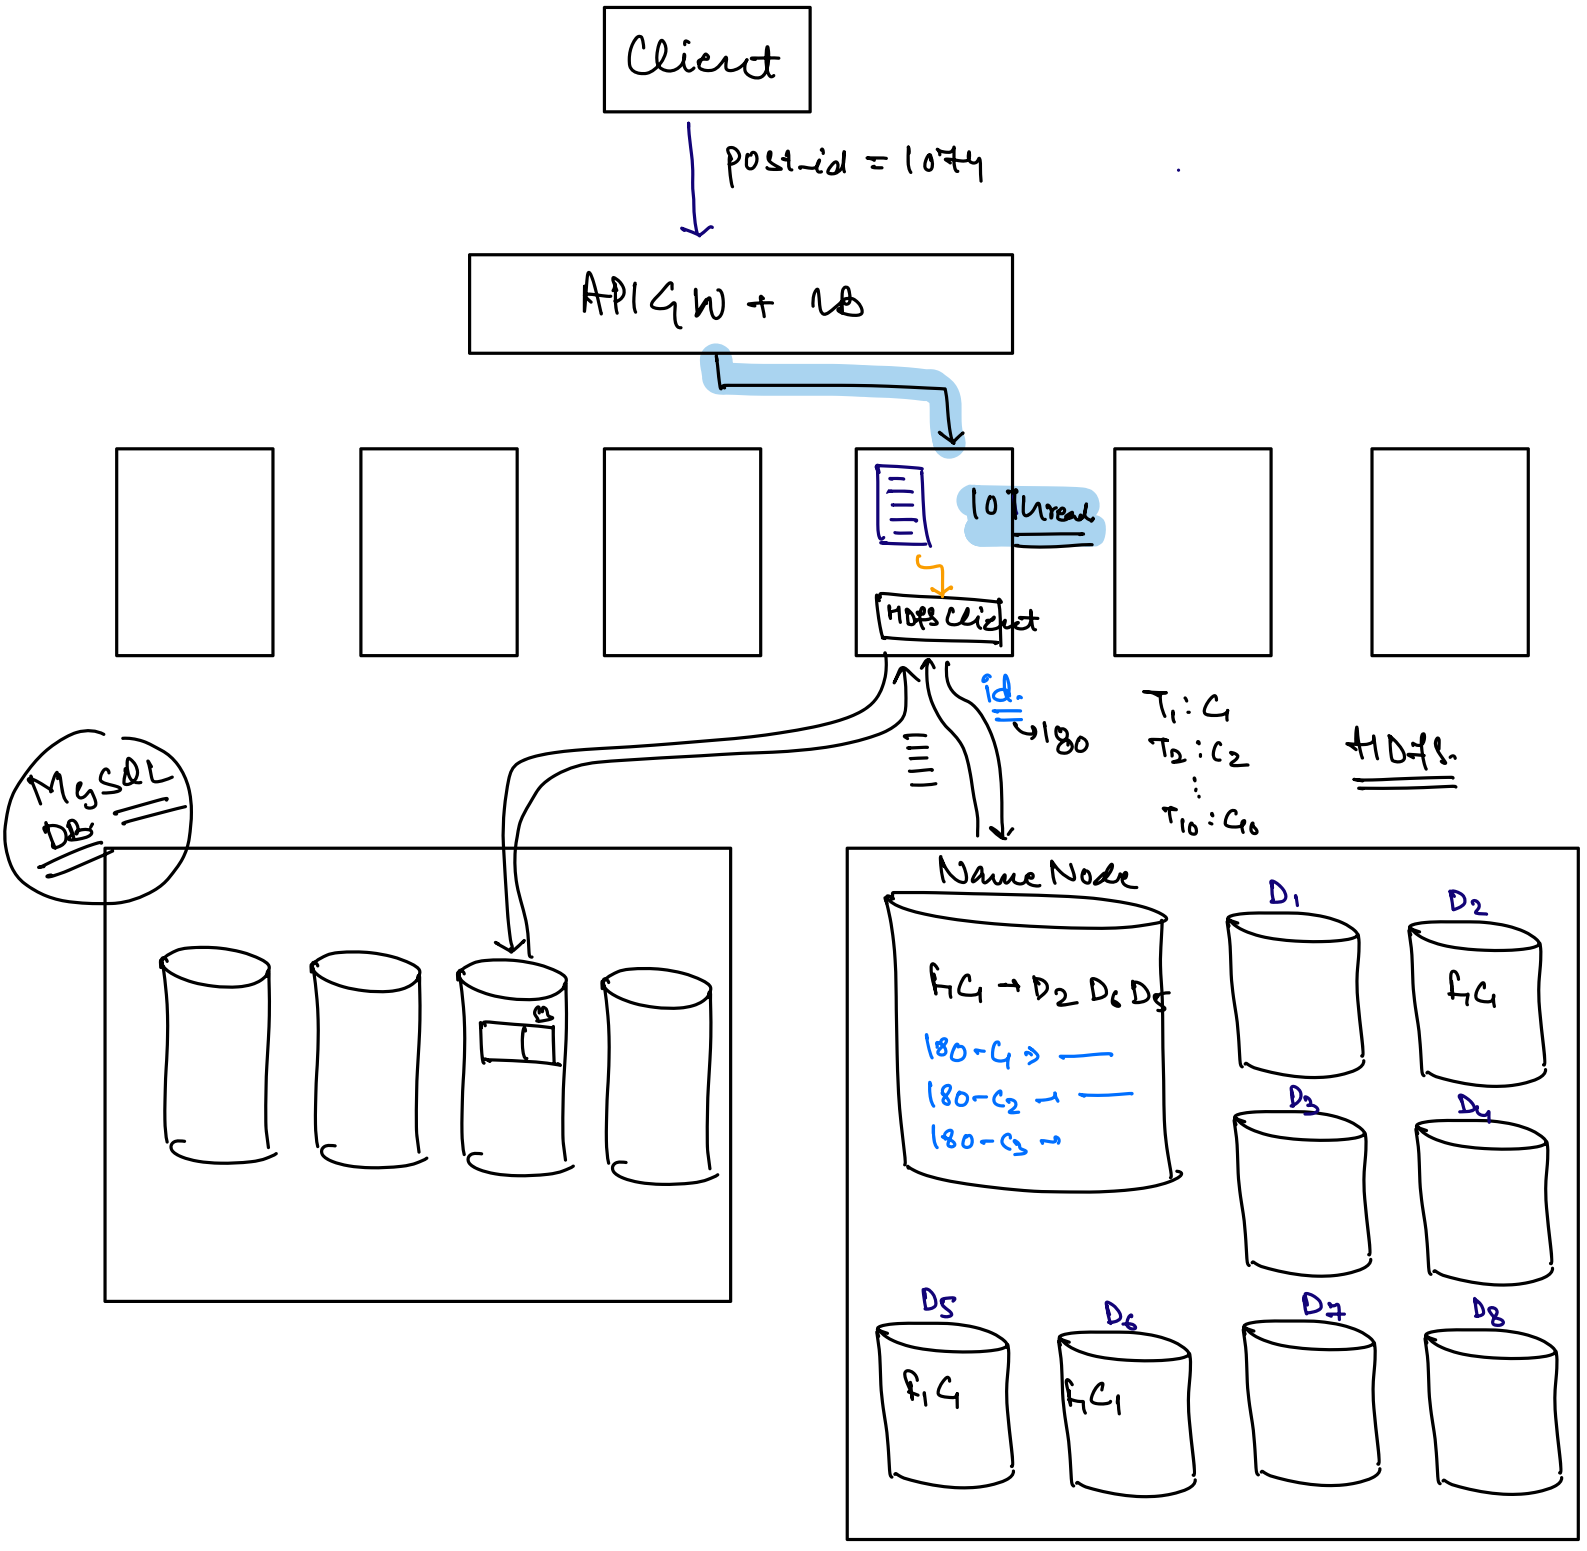
Client
post_id = 1074
API GW + LB
10 Thread
HDFS Client
id.
180
T1 : C1
T2 : C2
T10 : C10
HDFS.
MySQL
DB
Name Node
f1 C1 → D2 D6 D5
180-C1 →
180-C2 →
180-C3 →
D1
D2
f1 C1
D3
D4
D5
f1 C1
D6
f1 C1
D7
D8