Agenda.

⇒ NoSQL Internals.

↳ How NoSQL DBs stores data internally.

SQL DB.

↳ Structured Data

↳ fixed Schema.

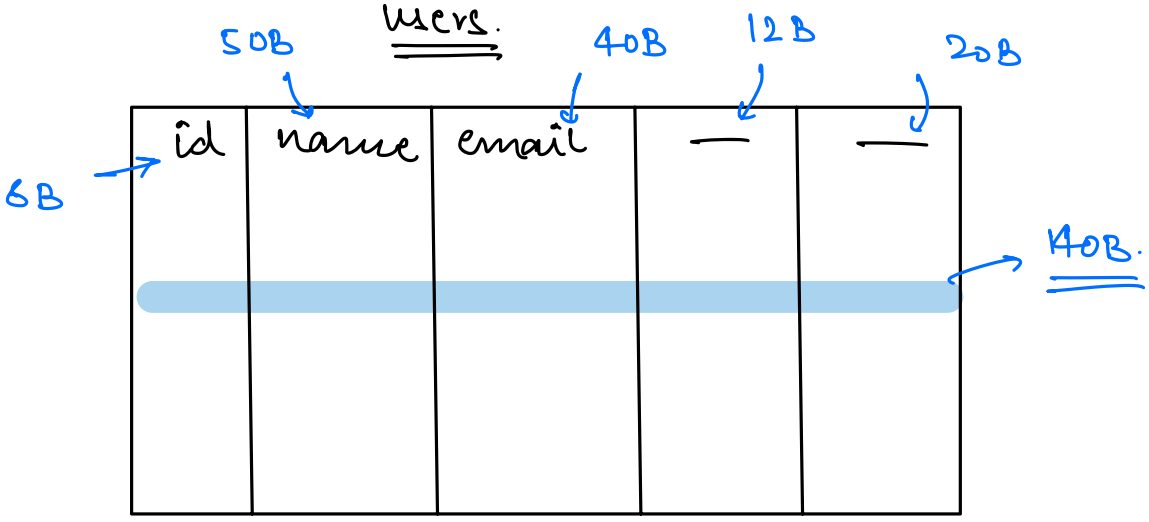

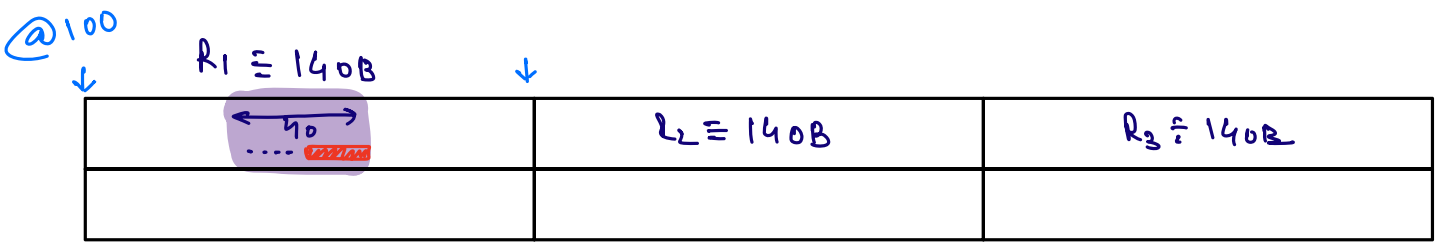

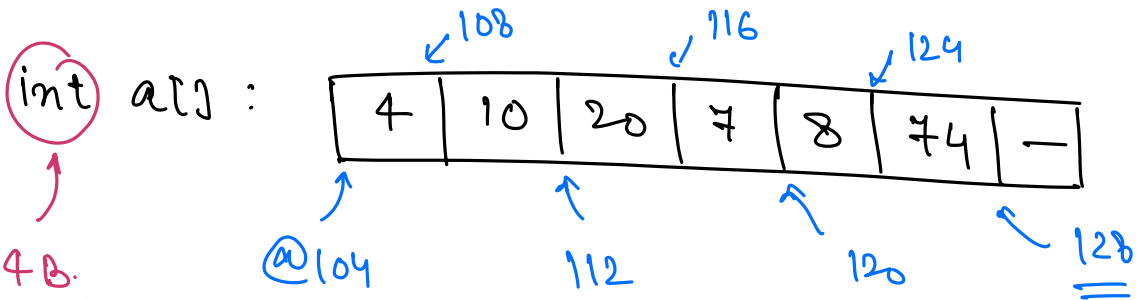

a[i] ↙ a[0] + i x 4

O(1)

NoSQL DB.

- Unstructured Data
- No fixed Schema.

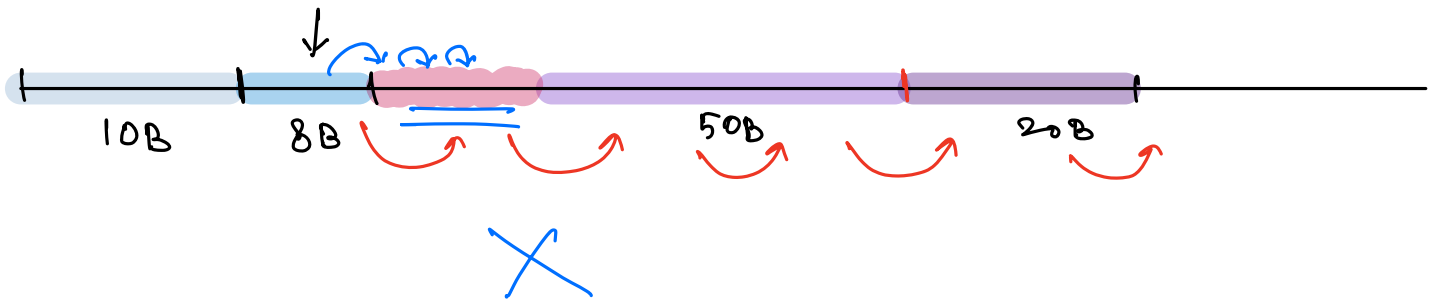

WAL

- Write Ahead log.
- Append Only file.

C
R
U
D
→ Write

WAL

A = 12
B = 40
C = 20
A = 24
D = 100
B = 41
E = 500
C = -1
A = 30
EOF →

K, V

WAL file

- Maintains the complete history of operations.
- Reconstruct our DB from WAL.
- Used to sync Replicas

Write op^n TC in WAL ⇒ O(1)

Read op^n TC in WAL ⇒ O(N) → # of write op^n

↳ Start reading the file from starting till the end and get the latest value of the key.

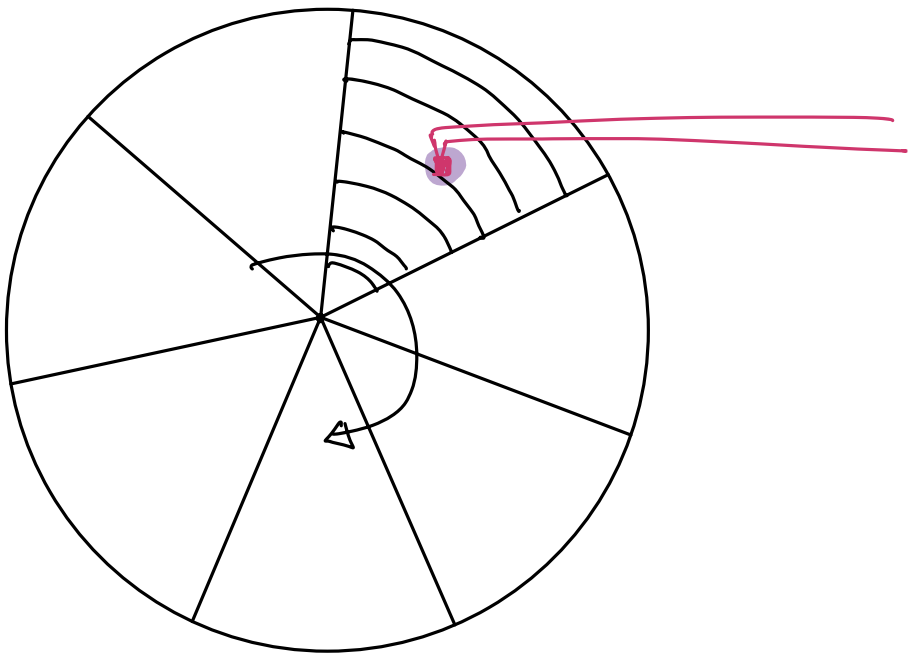

SQL DB

↳ Read TC : O(logN)

↳ Write TC : O(logN)

Approach #1

→ Only WAL.

Read TC : O(N)

Write TC : O(1)

Approach #2

→ WAL + HashMap.

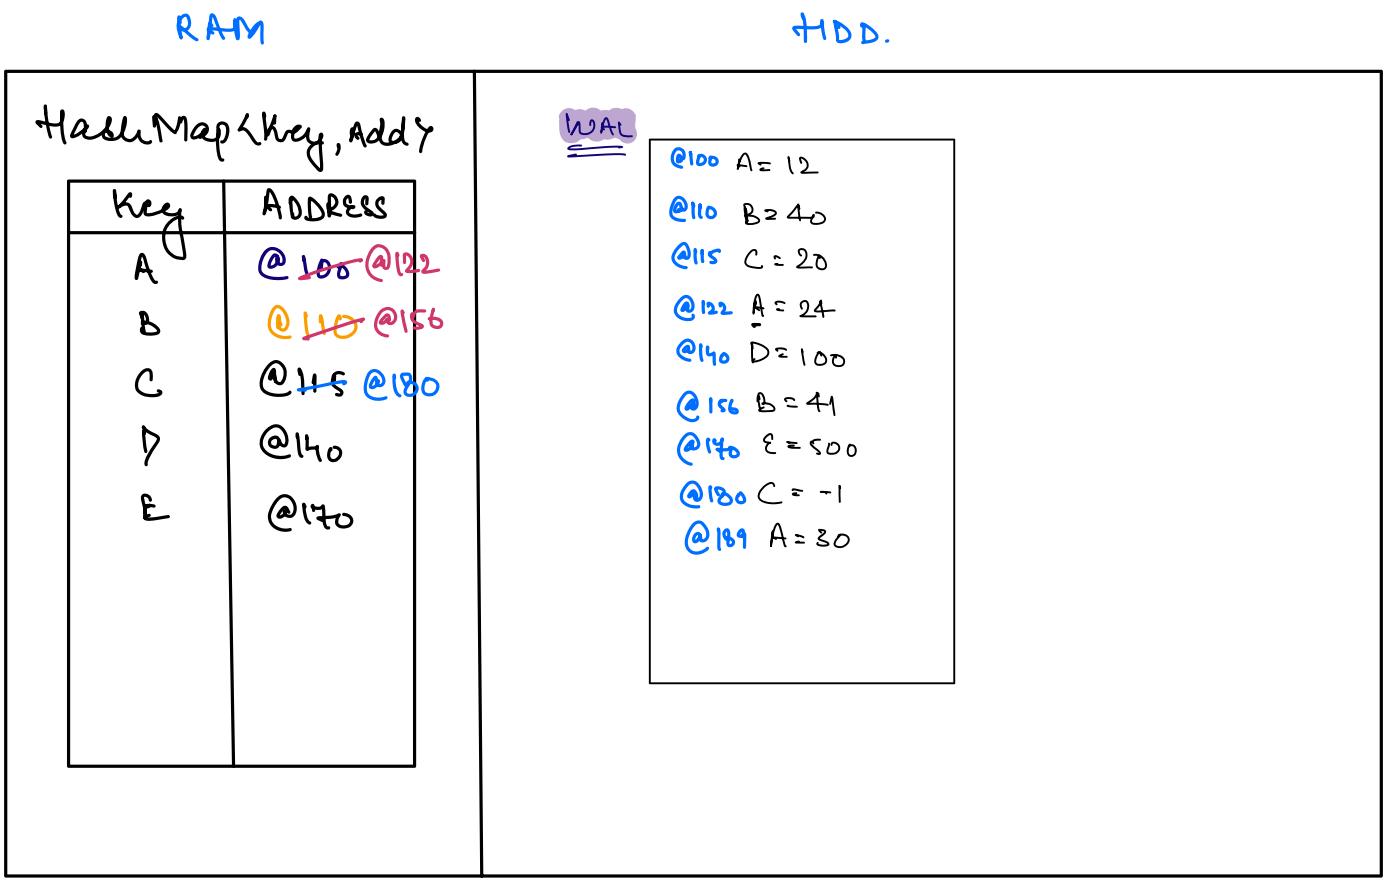

Write Opⁿ

↳ Append in WAL ⇒ O(1)

↳ Update HM ⇒ O(1)

TC of write ⇒ O(1)

Read Opⁿ ⇒ O(1)

↳ Get the address of the key from HM and read the value from this address in WAL

Cons.

↳ HM is present in the RAM (which is volatile storage), in case our m/c restarts, we'll have to rebuild the complete HM from scratch.

Size of WAL = 10 TB

Size of 1 entry = 10 B.

No of entries = $\frac{10\,TB}{10\,B} = \frac{10 \times 10^{12}\,B}{10\,B}$

= 10^{12} entries.

Assumptions: Unique keys = $\frac{10^{12}}{10} = 10^{11}$

Size of 1 K,V pair in HM ≈ 16 B.

Size of the HM = $10^{11} \times 16$ B.

= 1.6×10^{12} B.

= 1.6 TB.

⇒ Size of the HM can be huge, it will cost huge amount of money.

Background Script.

⇒ Every 1 hour, take the Data from HM and put it in the file in HDD & reset the HM.

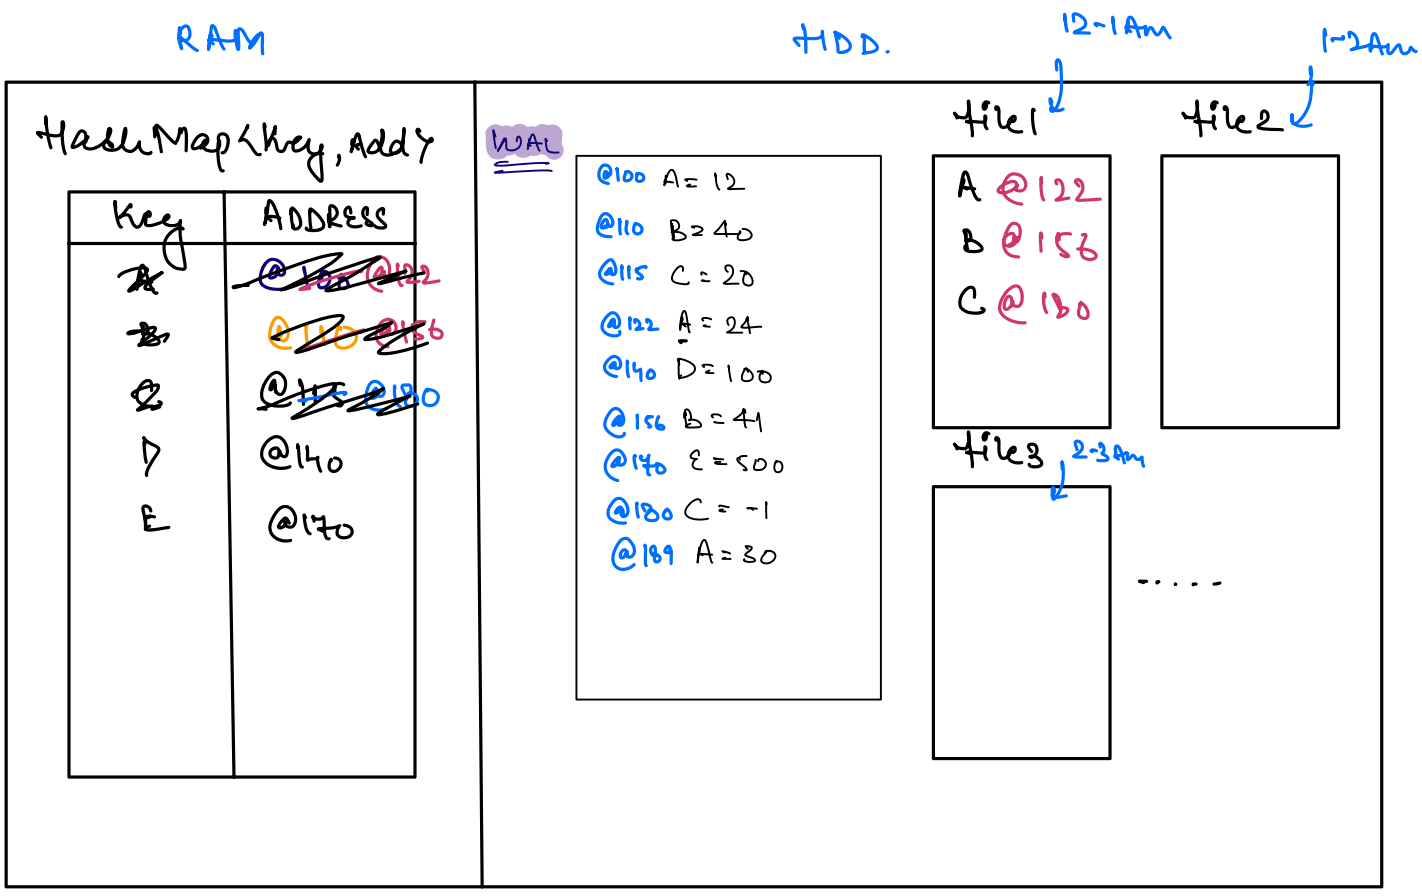

Write

- Write in WAL ⇒ O(1)
- Write in HM ⇒ O(1)

Read.

→ find the key in the Map

if found

Get the Address from HM & read from WAL.

if NOT found

Read all the files from latest to oldest (Max 24 files)

TC: O(1) + (# of files) × N

↓

HM

: (# of files) × N.

→ Iterate every hourly file linearly.

Summary.

	Read TC	Write TC
Only WAL	O(N)	O(1)
WAL + HM	O(1)	O(1)
WAL + HM + Hourly files	(# of files) × N.	O(1)

Binary Search

↳ Sorted Dataset + Equal size Data.

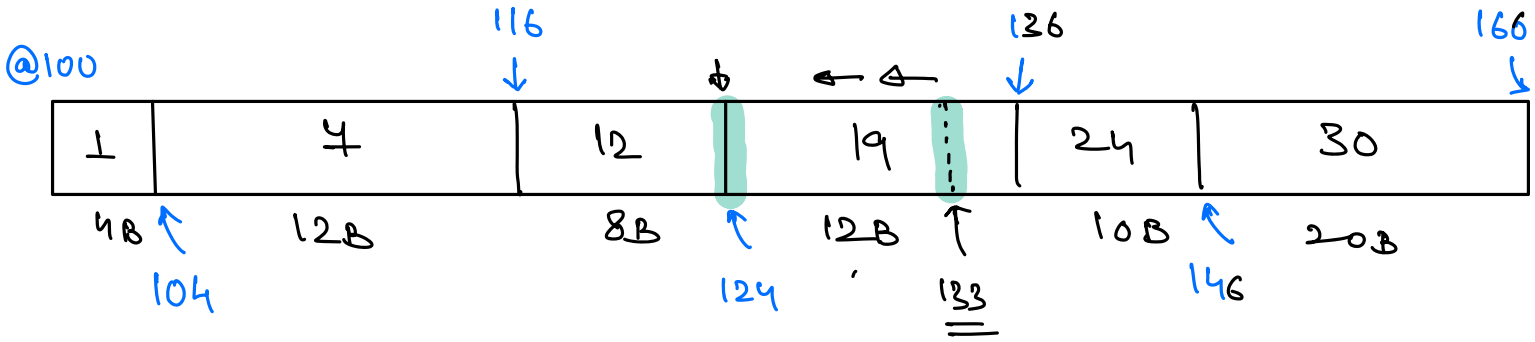

$$\frac{100 + 166}{2} = 133$$

⇒ Instead of HashMap, use Tree Map.

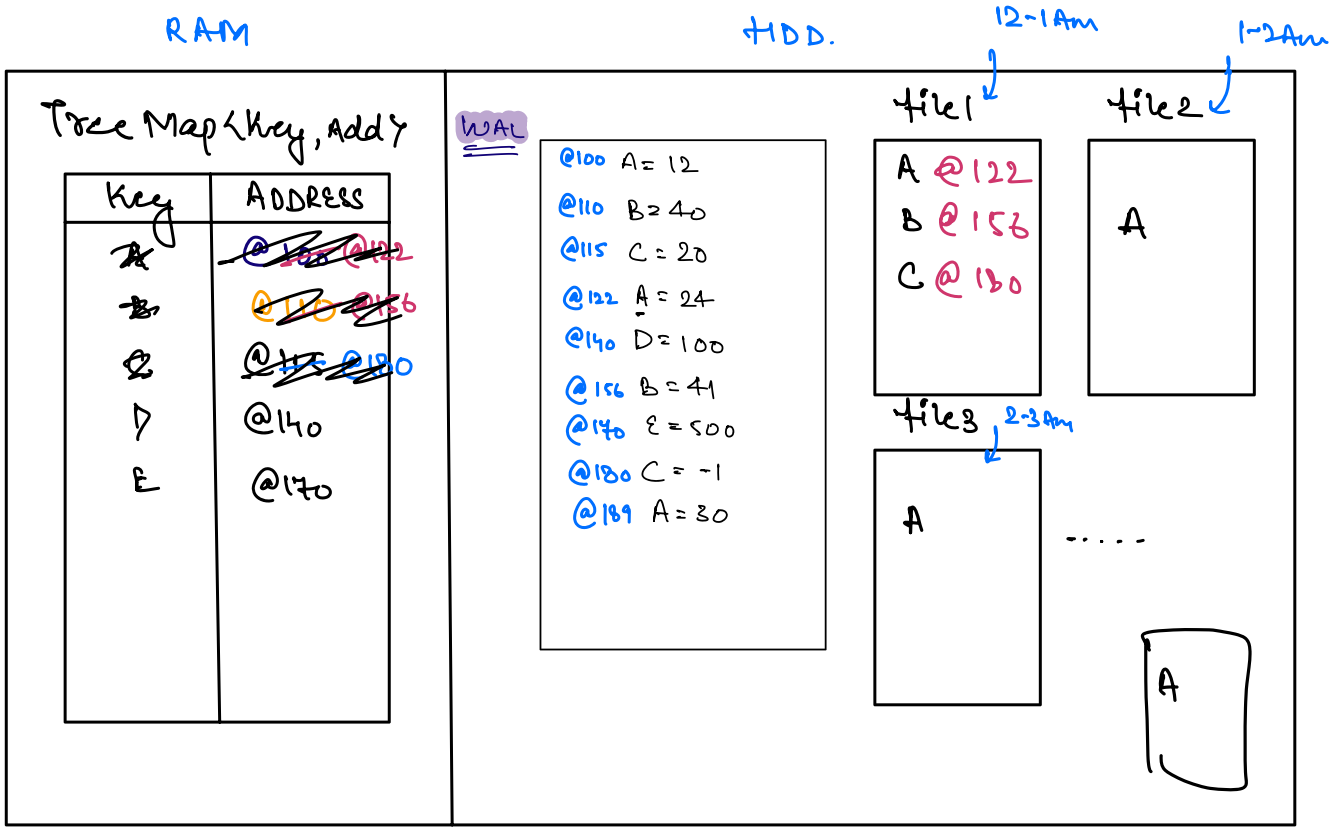

Write

- Write in WAL ⇒ O(1)
- Write in tree Map ⇒ O(logN)

} O(logN)

Read.

- find the key in the TreeMap
  - if found
  - if NOT found

Read TC : O(logN) + (# of files) × N

Summary.

	Read TC	Write TC
Only WAL	O(N)	O(1)
WAL + HM	O(1)	O(1)
WAL + HM + Hourly files	(# of files) × N	O(1)
WAL + TM + Hourly files	(# of files) × N	O(logN)

2 Sorted Array.

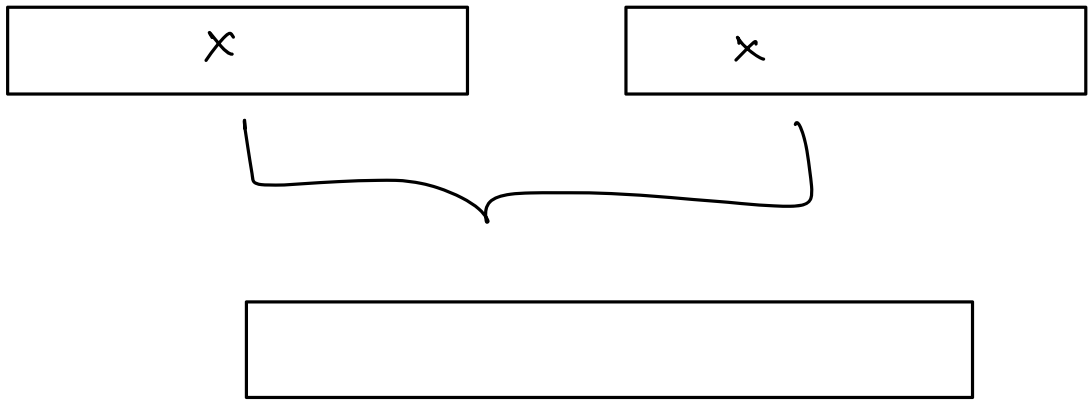

Background Script.

⇒ Every 24 hours, merge all the files (from file 1 to file 24) into 1 sorted file.

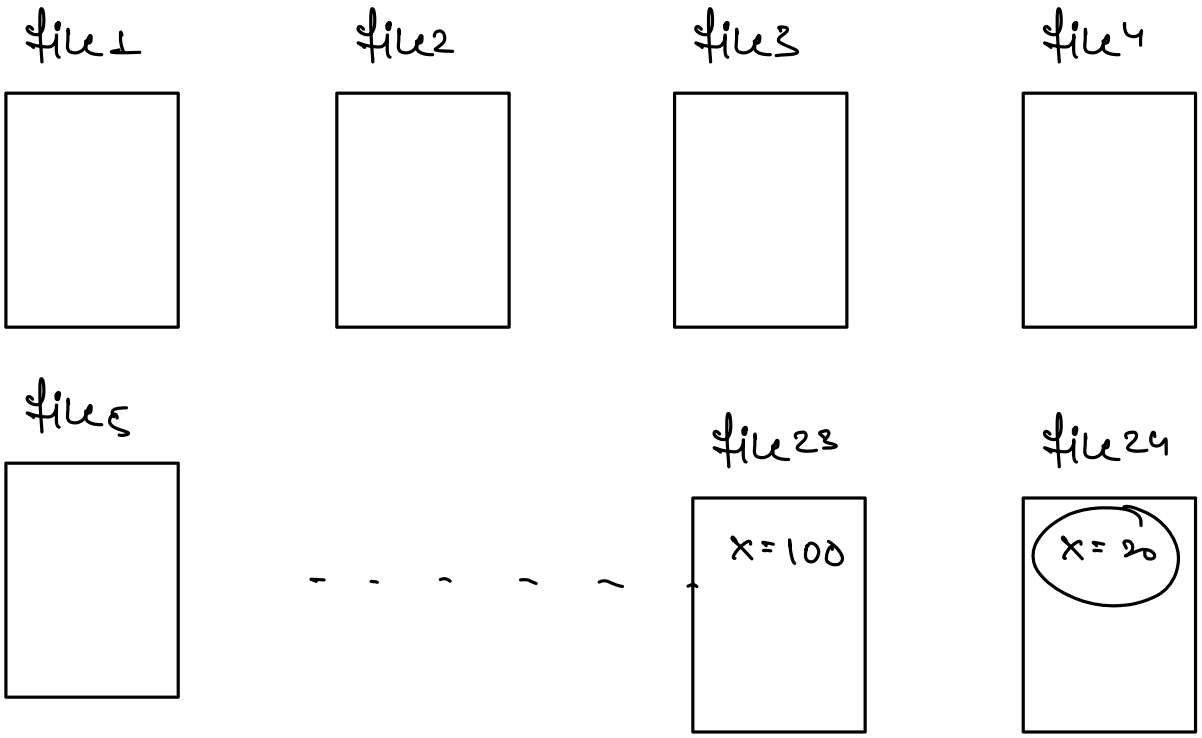

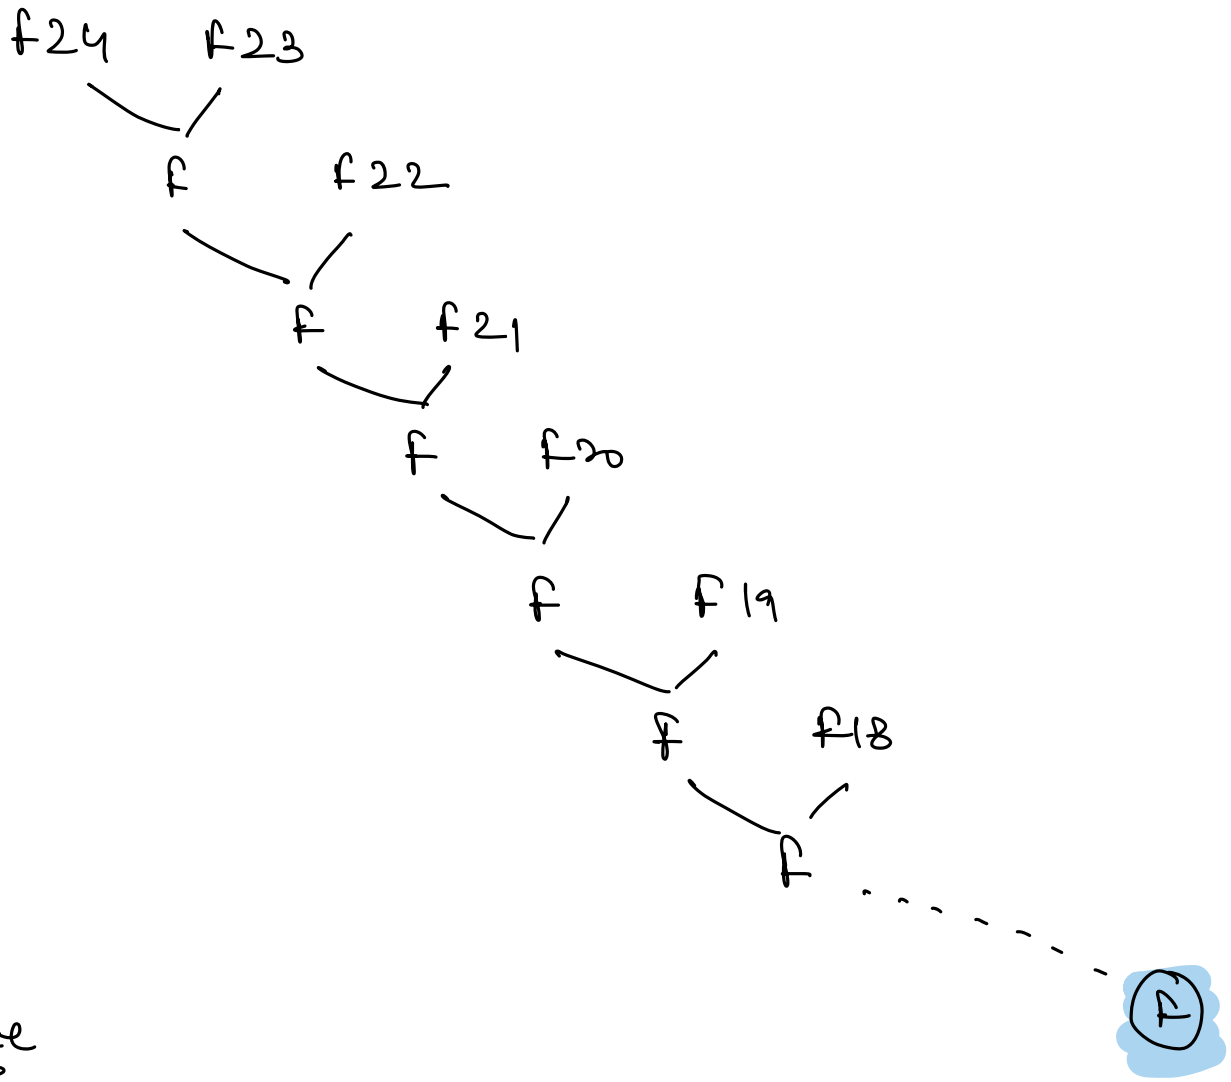

Write

- Write in WAL ⇒ O(1)
- Write in TreeMap ⇒ O(logN)

} O(logN)

Read.

- Read from TreeMap.
  - if found
    - Get from WAL
  - if NOT found
    - Iterate through all the files from latest to oldest.

Read TC : $O(\log N) + (\#\text{ of files}) \times N$.

Soln : Some how if we can apply Binary Search on the files, then we are done.

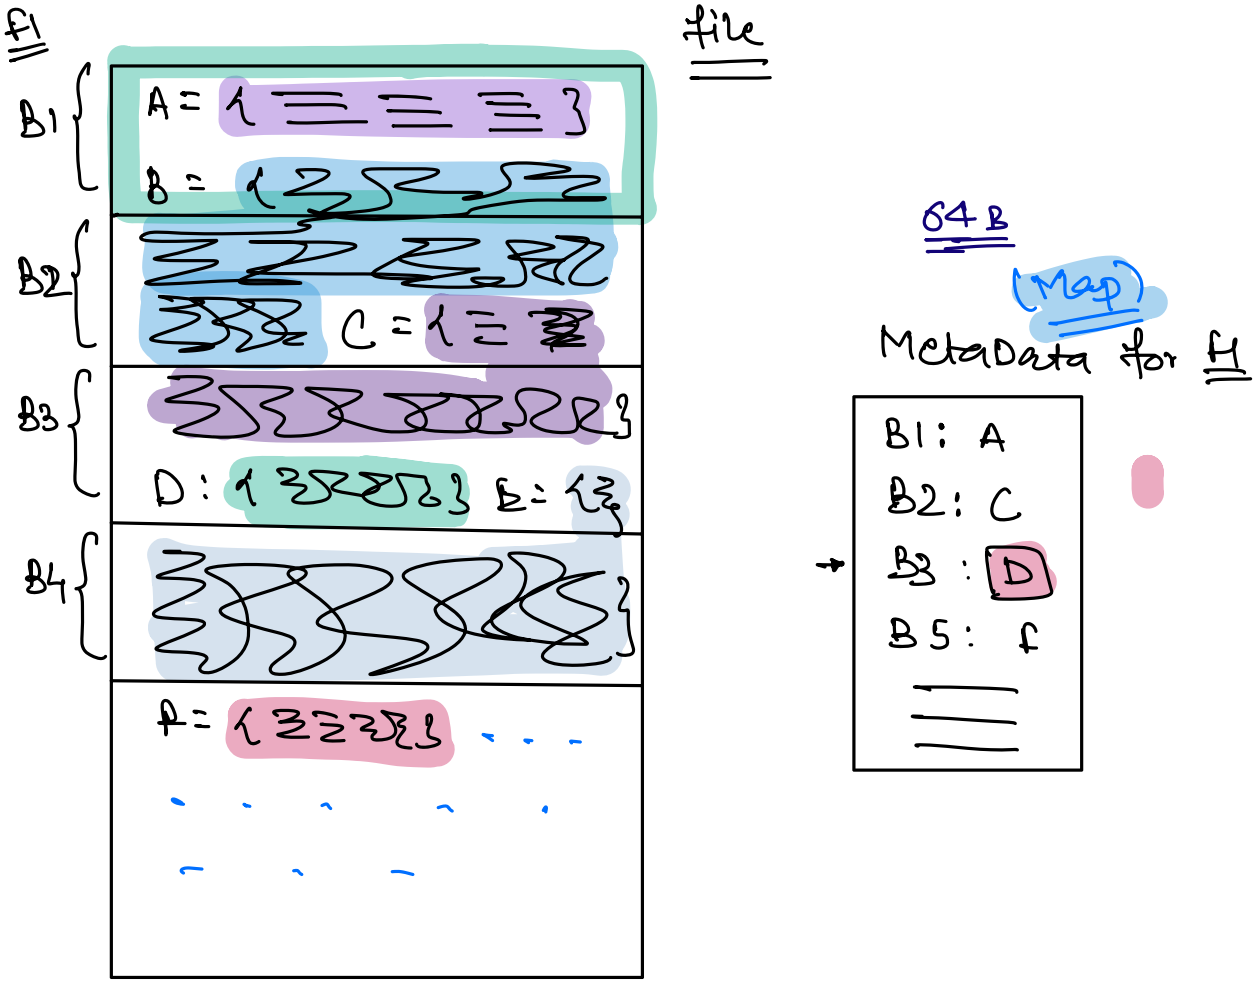

Final Soln

Write

- Write in WAL ⇒ O(1)
- Write in TreeMap ⇒ O(logN)

} O(logN)

Read.

- Read from TreeMap.
  - if found
    - Get from WAL
  - if NOT found
    - ⇒ Apply Binary search on the files from latest to oldest.

F24 → F23 → F22 ..... F1 → (F)

logN logN logN logN logN

Read TC ⇒ logN + (# of files) × logN.

= logN.

Log file (WAL) + Sorted Strings + Merge + Trees (Tree Maps)

$\equiv$ LSM Trees