Messaging Queues.

→ Asynchronous.

→ Messages / Events are placed in the Queue by the Producers without waiting for the Subscribers to process them.

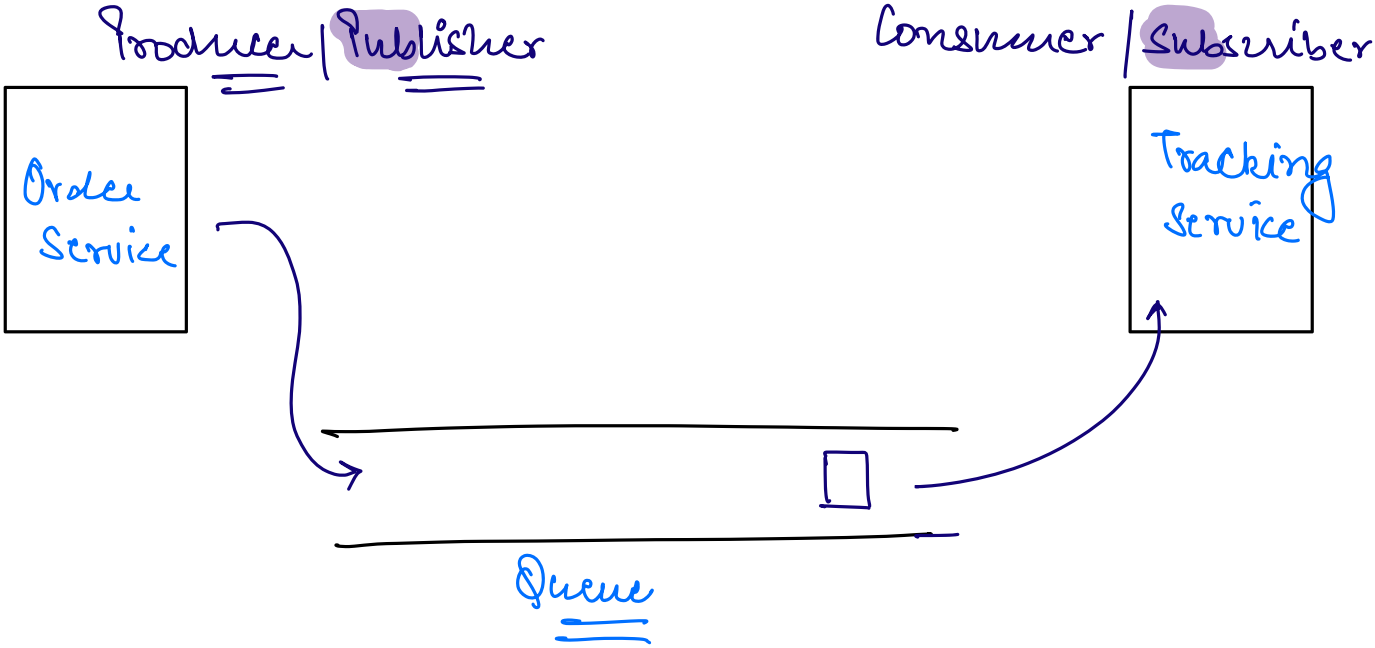

Advantages of MQ.

1) Decoupling.

→ Scalability: Producers & Consumers can be scaled independently.

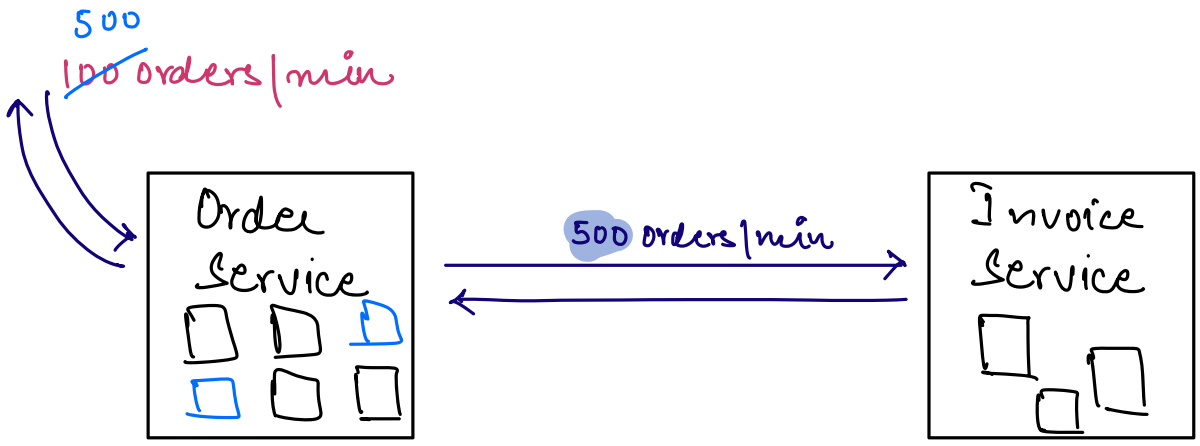

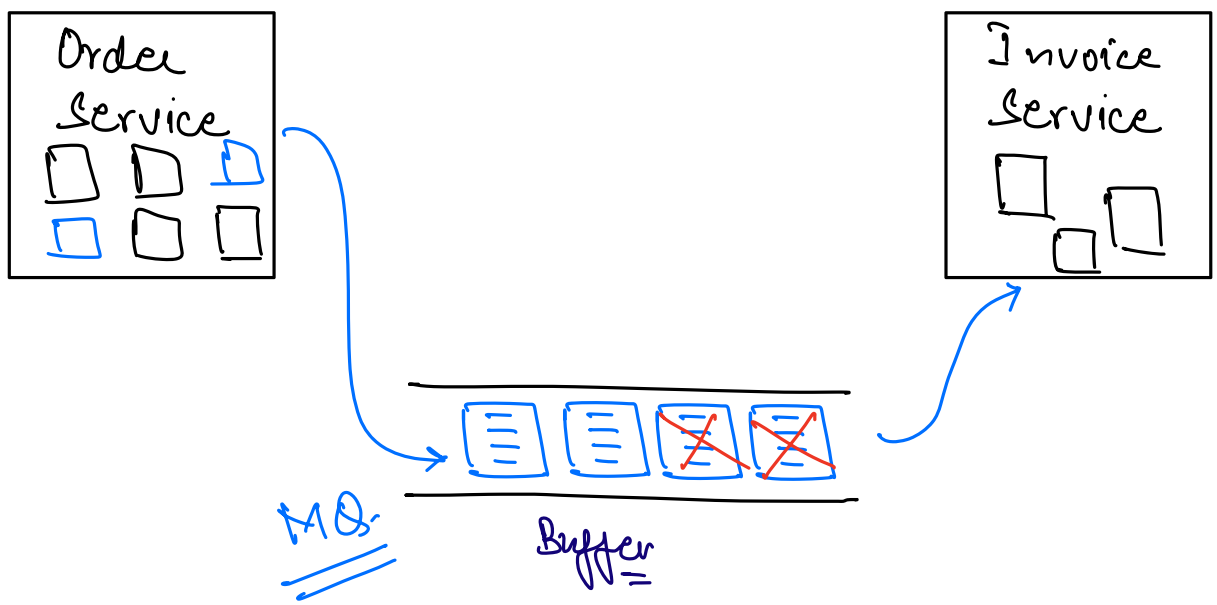

→ Fault Tolerance

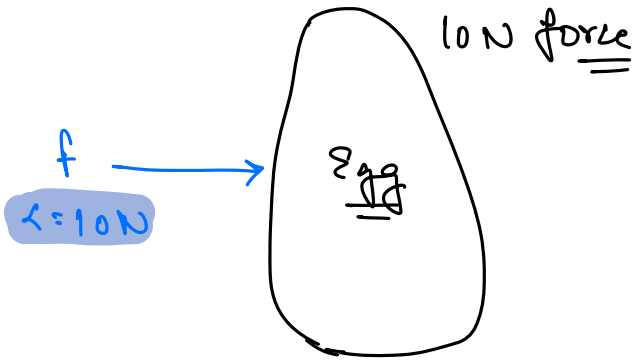

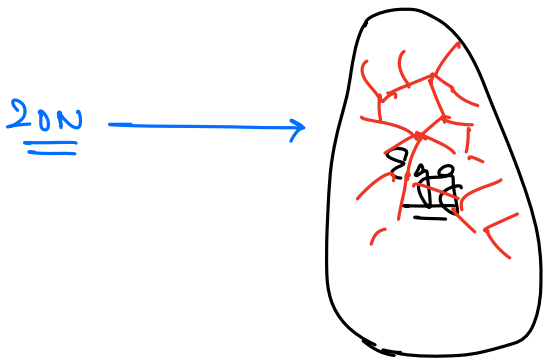

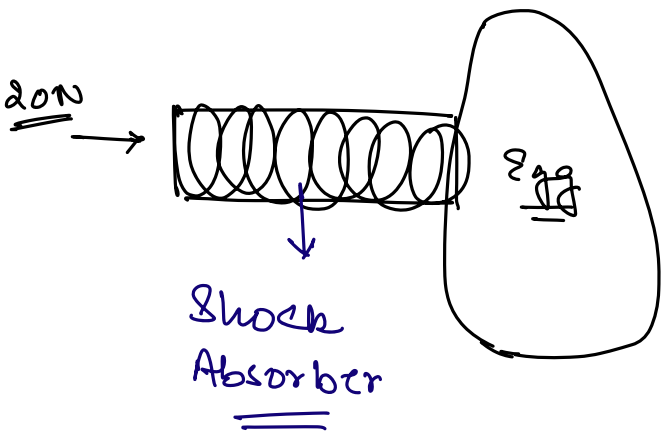

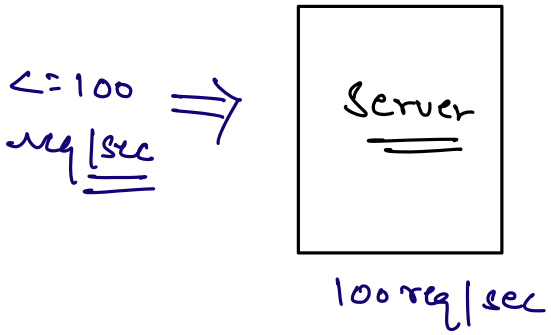

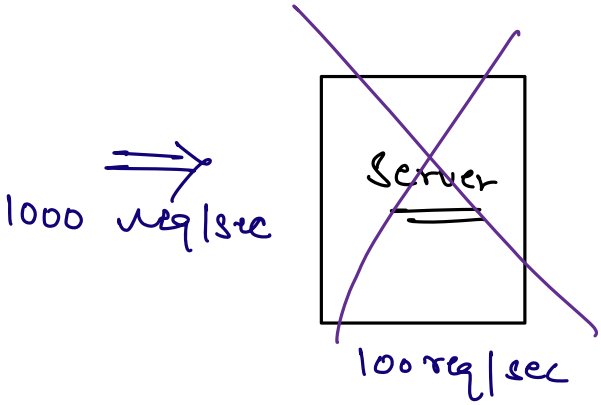

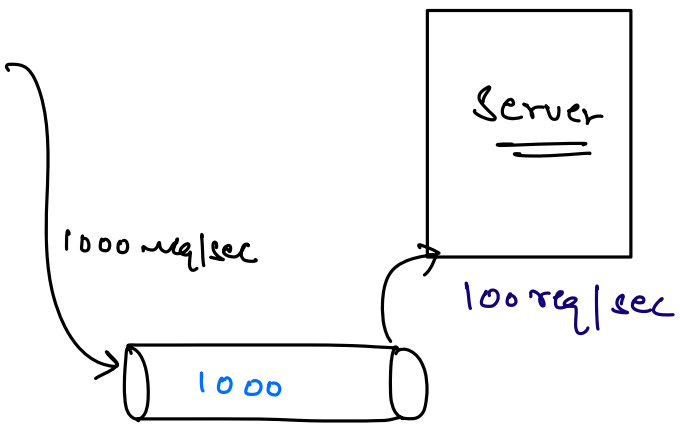

⇒ Buffering & Persistence

Use Cases.

1) Communication b/w Microservices.
2) Async processing | Task scheduling.

Cons of MQ.

→ Can't perform sync processing.
→ There can be a High latency.
→ Needs additional infrastructure.
→ High N/w Overhead.

Examples of MQ

→ Kafka.
→ Rabbit MQ
→ Amazon SQS
  ↳ Managed Infra.

1. Open Source: built at Linkedin, now part of Apache
2. High Throughput: a cluster can handle up to 10 million msg/s (trillion/day)
3. Ultra Low Latency: claims < 2ms latency from message insertion to consumption
4. Horizontal Scalability: automatically manages partitions across multiple servers (brokers)
1000s of brokers (servers), 100,000s of partitions, petabytes of data
5. Fault Tolerant: automatically replicates messages across multiple servers
6. High Availability: can deploy clusters across availability zones
7. Persistent: saves the messages on disk to ensure they're not lost. Multiple consumer groups can consume the same message.
8. Message Ordering: provides strict* message ordering (*within each partition only)

⇒ msg | event = { "key" : ——— "value" : ——— ——— ——— ——— }

≤

⇒ The size of the msg | event should be limited to few KBs

Producers & Consumers.

Writes to kafka.

reads from kafka.

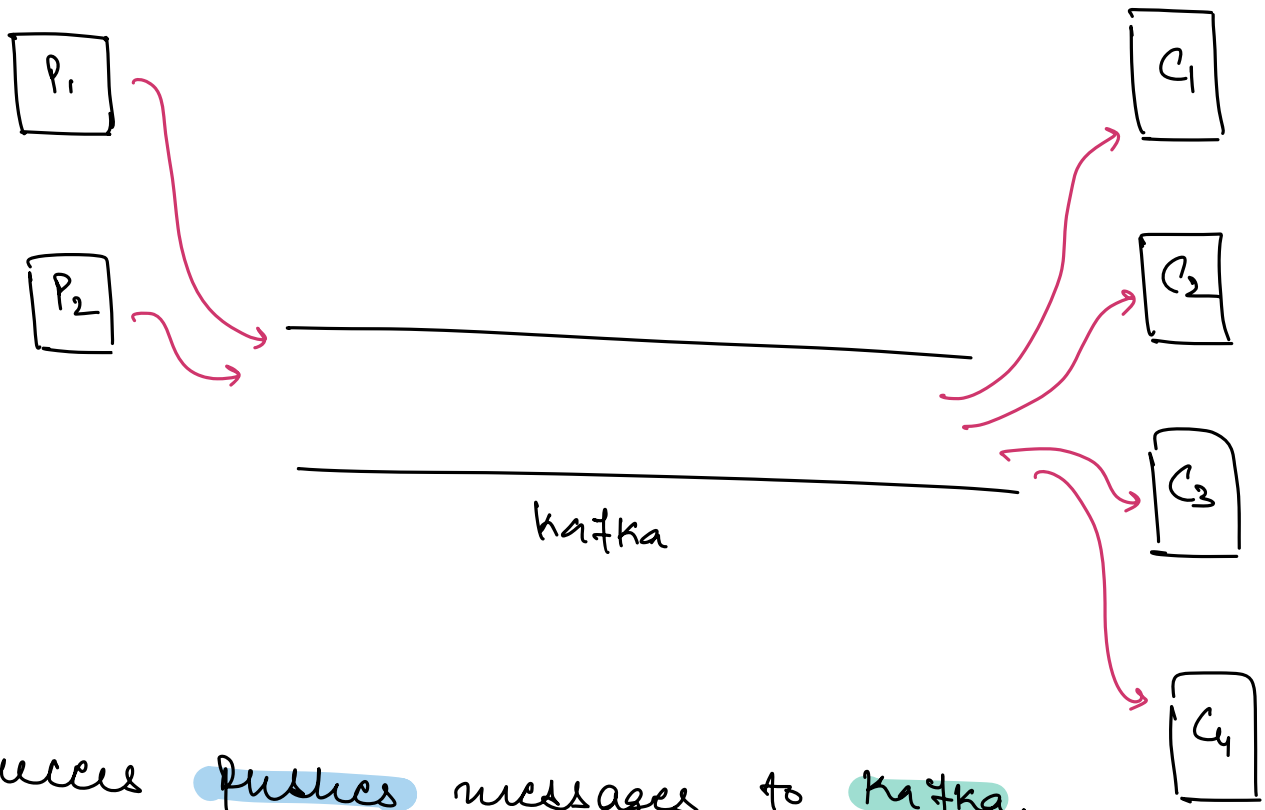

Producers publish messages to kafka.

Consumers pulls messages from kafka.

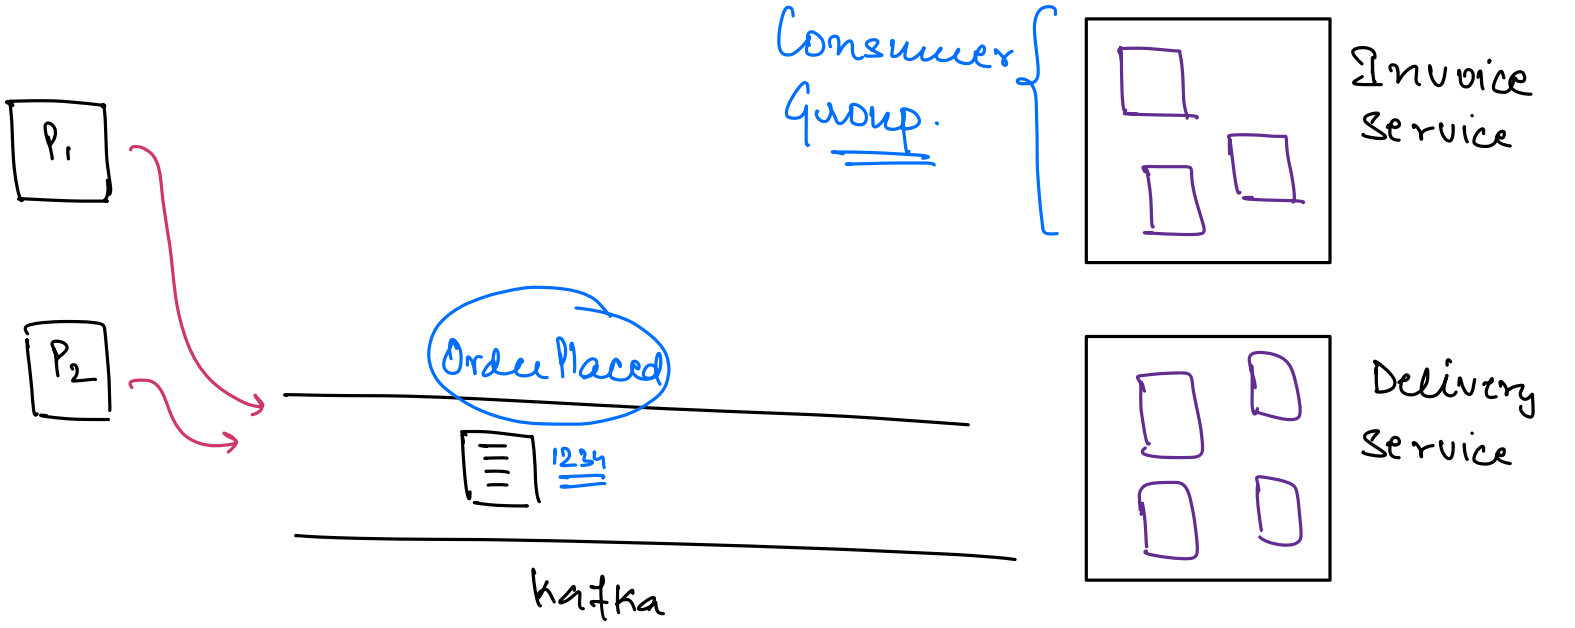

⇒ An event can be consumed by multipe Consumer Groups, but within the group it should be consumed only by one Consumer.

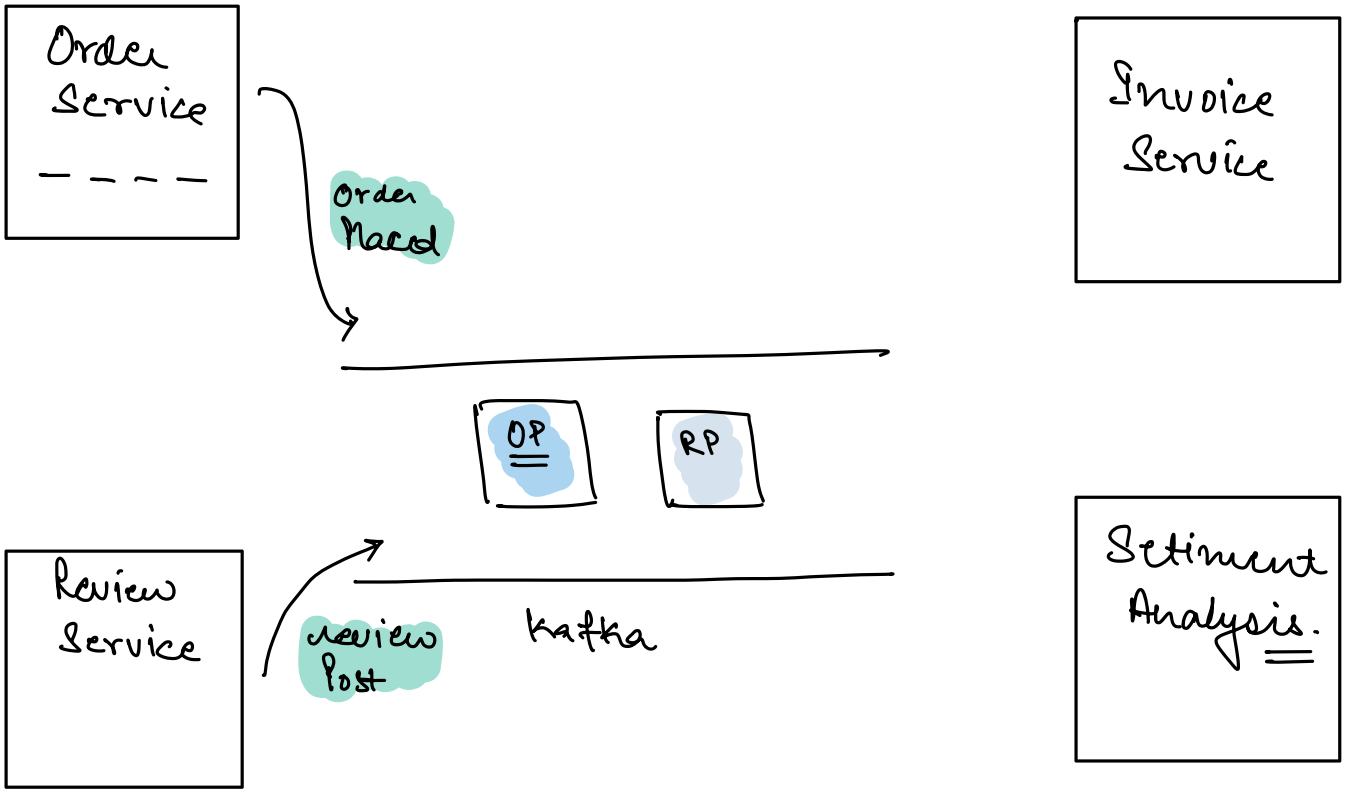

⇒ Different consumer groups can choose to consumes specific category of messages.

⇒ Categorize messages into Topics.

⇒ A msg can belong to exactly one Topic.

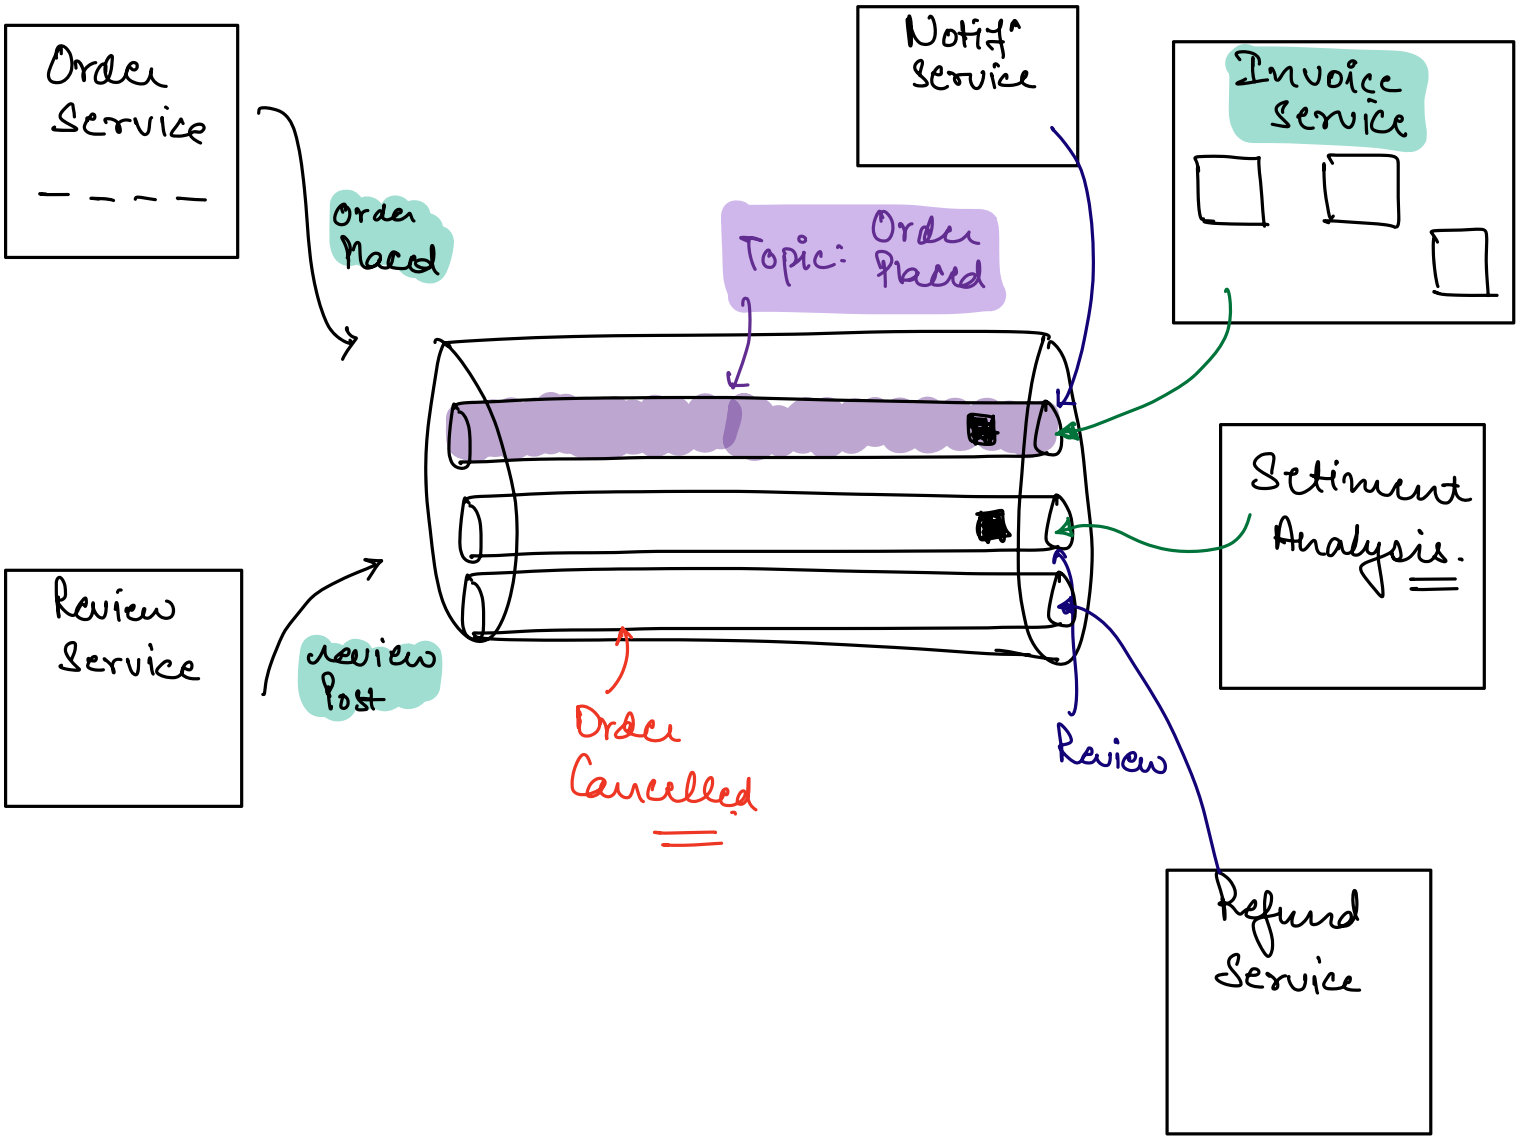

Topics in Kafka are always multi-producer and multi-subscriber: a topic can have zero, one, or many producers that write events to it, as well as zero, one, or many consumers that subscribe to these events

PARTITIONS

→ A single server is NOT enough to store all the messages in a topic.

Ex: FB Messages.

→ As the topic is too large to store in a single server, we can partition the topic across multiple servers.

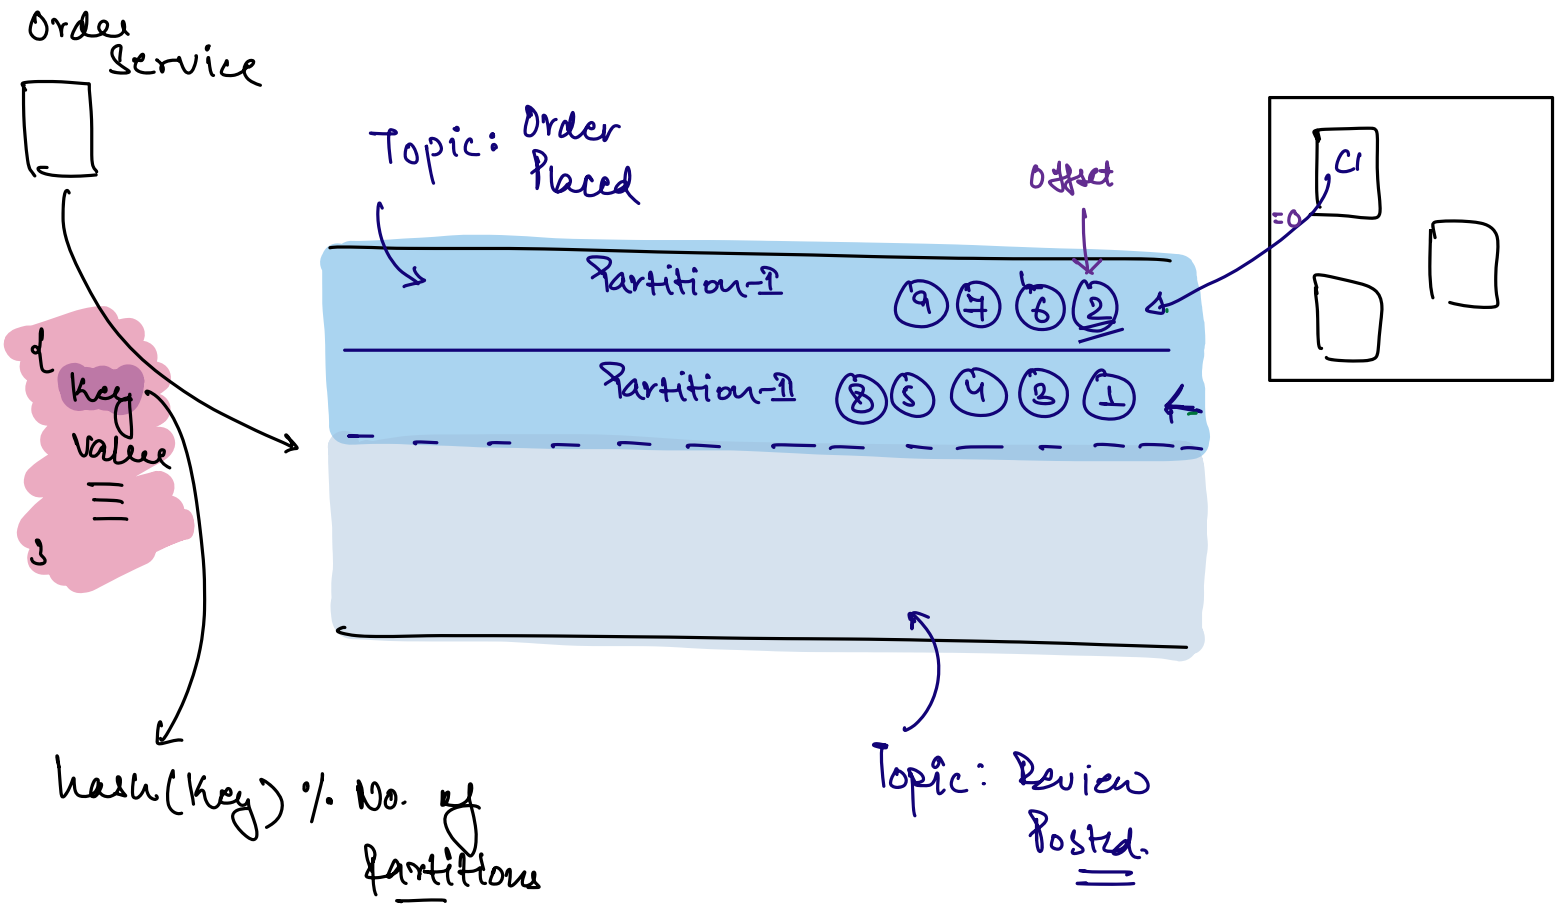

⇒ # of partitions = # of Consumers

↳ No consumer will be sitting idle.

⇒ # of partitions < # of Consumers

↳ few consumers will be sitting idle.

⇒ # of partitions > # of Consumers

⇒ We must ensure the # partitions >= # Consumers

Note: 1 partition is assigned to exactly 1 consumer but 1 Consumer can get multiple partitions.

⇒ If we have more consumers than partitions, some consumers will remain idle

⇒ Since we might need more consumers in future, so we take good enough large # of partitions

⇒ Kafka ensures the strict ordering with the partition and NOT across partitions.

Scalability

- SHARDING
- REPLICATION

M/C ≡ Server

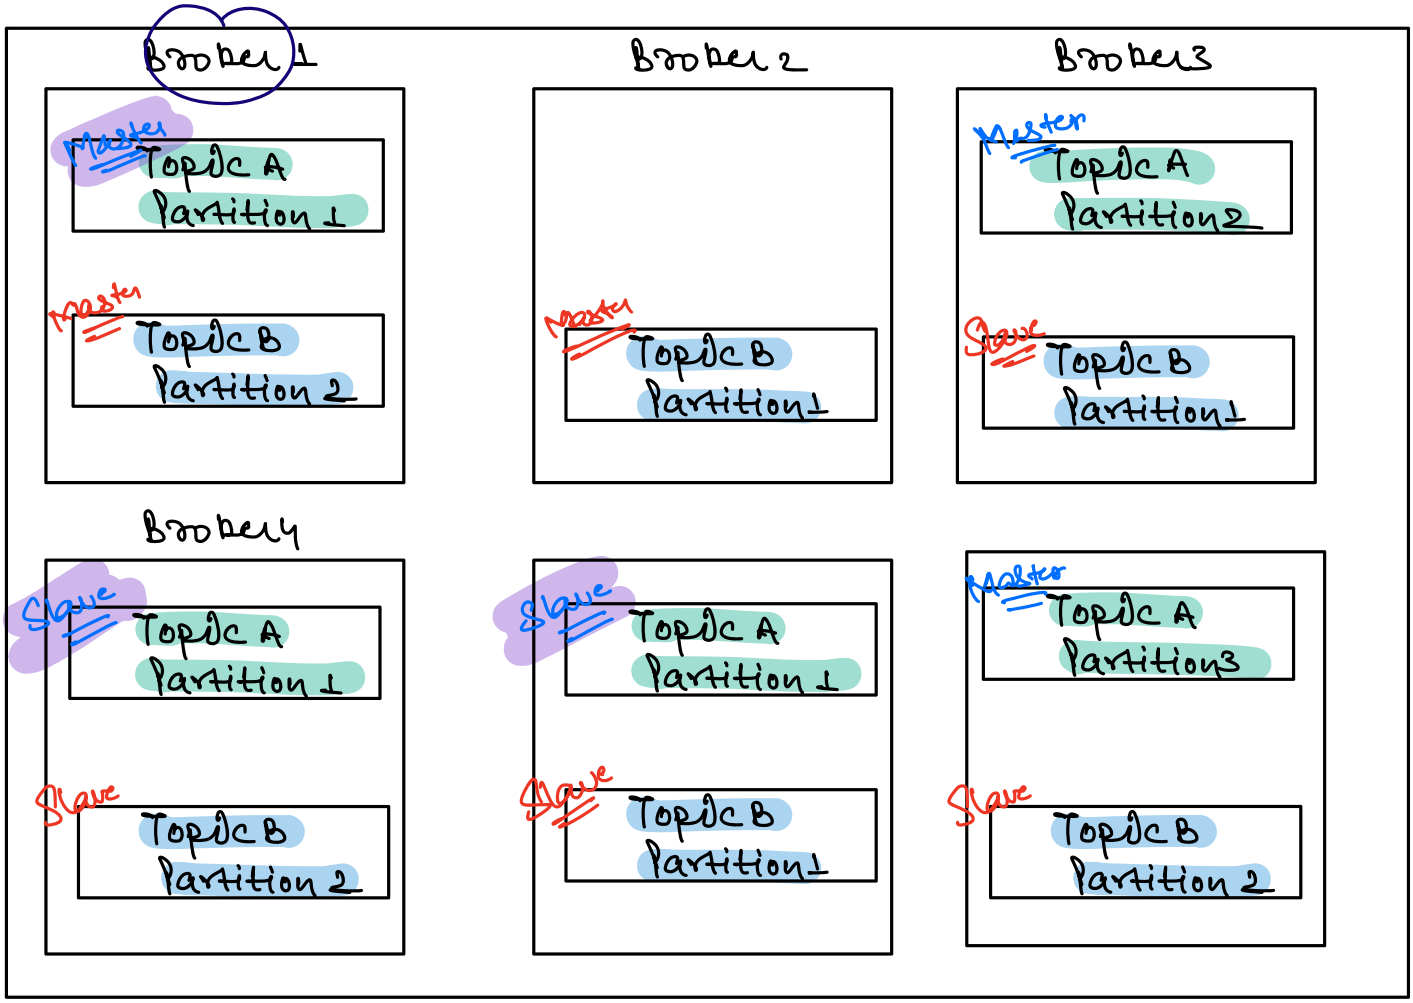

ZOOKEEPER